

Generative Models for Opponent Modeling in Poker

Kai Liu

Cornell University

k12298@cornell.edu

Xingwei Li

Cornell University

xl988@cornell.edu

December 14, 2025

Abstract

This project investigates generative approaches to opponent modeling in poker from two complementary perspectives. First, we study persona-conditioned large language models (LLMs) as controllable generators of diverse, human-interpretable opponent behaviors. Second, we embed a learned generative opponent model directly into a reinforcement learning (RL) framework, allowing an agent to adapt its policy based on predicted opponent actions. Together, these two components demonstrate how generative models can both simulate strategic diversity and inform adaptive decision-making in imperfect-information games.

1. Introduction

Poker is a classic testbed for strategic decision-making under uncertainty. As an imperfect-information game, it requires players to infer hidden information and adapt to opponents whose strategies are only partially observable. A central challenge in this setting is *opponent modeling*, which plays a key role in effective decision-making for both human players and artificial agents.

Recent poker AIs based on reinforcement learning, such as DeepStack (Heimrich and Silver, 2016) and Libratus (Brown and Sandholm, 2018), have achieved superhuman performance through self-play and equilibrium reasoning. While these methods produce strong and theoretically sound strategies, they are primarily designed to be unexploitable and often fail to adapt to human-like opponents with diverse behavioral patterns.

Generative Artificial Intelligence offers a promising alternative. In poker, generative models can be used to simulate and predict opponents with different playing styles, enabling training and evaluation beyond standard self-play. In this project, we investigate the use of generative models for opponent modeling in poker.

2. Background Information

2.1. Poker Rules and Key Concepts

Texas Hold’em poker is played with two private hole cards per player and five shared community cards, from which each player forms the best five-card hand. A hand proceeds through four betting rounds: preflop, flop, turn, and river. At each round, players may choose among

actions such as fold, check, call, bet, or raise. Player position is strategically important, with late positions generally having an informational advantage.

To characterize player behavior, poker analytics commonly use summary statistics. These include VPIP (Voluntarily Put Money in Pot), PFR (Preflop Raise frequency), aggression frequency (AFq), and showdown-related metrics such as WTSD (Went to Showdown) and W\$SD (Won Money at Showdown). Such statistics provide compact representations of strategic tendencies and are widely used in opponent analysis.

Leduc Hold’em is a simplified variant of Texas Hold’em that retains the core imperfect-information structure of poker while substantially reducing state and action complexity. The game is played with a small deck consisting of two copies of each rank, each player receives a single private card, and exactly one public card is revealed during the hand. Betting proceeds over two rounds—before and after the public card is revealed—with a limited action space and fixed bet sizes. Despite its simplicity, Leduc Hold’em captures key strategic elements such as bluffing, value betting, and belief updating, making it a widely used benchmark for studying reinforcement learning and opponent modeling in poker-like environments.

2.2. Generative Models in Sequential Decision Problems

Generative models provide a principled way to model complex distributions and sequential data. Techniques such as GANs (Goodfellow et al., 2014), diffusion models (Ho and Salimans, 2020), and transformers (Vaswani et al., 2017) have demonstrated strong performance in generating realistic trajectories across a variety of domains.

In reinforcement learning, generative approaches have been extended to imitation learning and policy modeling. Methods such as Generative Adversarial Imitation Learning (Ho and Ermon, 2016) and Decision Transformers (Chen et al., 2021) model entire action sequences, suggesting that generative architectures can capture decision-making behavior over time.

2.3. Generative Opponent Modeling

Generative opponent modeling extends traditional opponent modeling by representing opponents as stochastic generators of behavior rather than fixed or deterministic predictors. Instead of outputting a single most likely action, generative models capture distributions over actions or strategies, allowing agents to reason about behavioral uncertainty, variability, and non-stationarity in opponent play.

In the context of reinforcement learning, generative opponent models can be learned from interaction data and used as predictive components within the agent’s decision-making process. Neural network-based opponent predictors, for example, can estimate an opponent’s next action or action distribution conditioned on the current state and interaction history. (Ekmekci and Sırin, 2013) When embedded into the agent’s state representation or value estimation, such predictive models enable online adaptation to exploitable or non-equilibrium behaviors, going beyond equilibrium-focused self-play.

Recent work using large language models (LLMs) in multi-agent games further demonstrates that generative models can simulate rich, context-aware opponent behaviors through natural-language conditioning (Project, 2024). These LLM-based simulators complement learned neu-

ral predictors by providing diverse, human-interpretable opponent strategies for evaluation and training. Together, simulation-based generative opponents and prediction-based opponent models form a unified framework for improving robustness and adaptability in reinforcement learning agents facing imperfect-information games such as poker.

3. Research Questions and Approaches

3.1. Research Questions

1. How can generative models capture the strategic diversity of real poker opponents?
2. Can neural network-based architectures generate action sequences consistent with game dynamics?
3. How can generated opponents be used to improve RL policy performance?

3.2. Approach

In this work, we first adapt a persona-conditioned LLM-based framework to study opponent modeling in poker. The LLM is prompted with natural-language persona descriptions corresponding to common poker archetypes and used to generate opponent actions in a controlled heads-up No-Limit Hold’em environment. Behavioral statistics and game outcomes against a fixed hero strategy are then analyzed to evaluate strategic consistency and separability across personas.

In addition to this approach, we consider an alternative modeling direction based on learned generative policies. In this setting, we build a generative model using multilayer perceptrons (MLPs). The model can be trained synchronically with a reinforcement learning agent to predict opponent actions conditioned on partial game trajectories. Such models offer a complementary, data-driven representation of opponent behavior that does not rely on natural-language conditioning.

Finally, generated opponent actions in the previous part are incorporated into reinforcement learning pipelines. By complementing reinforcement training with generated opponent behaviors, we analyze whether it is possible to improve performance compared to the baseline version without opponent modeling.

4. LLM-based Opponent Modeling

4.1. Overview and Research Question

Opponent modeling in poker seeks to characterize an opponent’s behavioral tendencies and leverage them to predict future actions. Recent advances in large language models (LLMs) suggest that such models can internalize high-level behavioral instructions through natural-language prompts and generate decision sequences that resemble human strategic styles.

This observation motivates the following research question: *Can persona-conditioned LLM prompts reliably produce strategically distinct and human-interpretable poker behaviors?*

To study this question, we construct a controlled two-player heads-up No-Limit Hold’em environment, in which the opponent (villain) is played by an LLM-based policy. The LLM is conditioned on one of several predefined persona descriptions representing common poker archetypes, including Solver, Nit, TAG, LAG, and Calling Station. The goal of this study is to evaluate whether natural-language persona prompts induce consistent, distinguishable, and human-like strategic behaviors in LLM-driven poker agents.

4.2. Environment and Simulation Setup

We design a simplified yet realistic two-player heads-up No-Limit Hold’em (NLH) simulation environment to study LLM-based opponent behavior in a controlled setting. The environment uses a standard 52-card deck and a five-card hand evaluator, with gameplay proceeding through the preflop, flop, turn, and river betting rounds. To reduce strategic complexity while preserving core decision-making structure, each betting round allows at most one bet followed by a response, preventing multi-raise sequences.

Each hand is treated as an independent episode, with player stacks reset at the start of every hand. The LLM-controlled opponent is provided with the full observable game state, including public cards and the complete action history. All actions and outcomes are logged at both the action level and hand-result level for subsequent behavioral and performance analysis.

The opposing player (hero) follows a fixed, rule-based strategy designed to approximate an aggressive, GTO-inspired regular. This strategy emphasizes higher preflop raising frequency with fewer flat calls, wider in-position defense, increased continuation betting and calibrated bluffing frequencies, and frequent river barrels while avoiding excessive out-of-position bluffs.

4.3. Persona Conditioning of the LLM

Opponent behaviors are induced by conditioning the LLM on natural-language persona descriptions that specify high-level strategic tendencies. We consider five common poker archetypes: Solver (balanced, GTO-like), Nit (tight and risk-averse), TAG (tight-aggressive and value-oriented), LAG (loose-aggressive and pressure-heavy), and Calling Station (loose-passive with minimal bluffing). These personas are provided as system-level prompts and remain fixed throughout each simulation run.

At each decision point, the LLM receives the assigned persona together with the full observable game state, including private cards, public board cards, pot size, stack sizes, current betting street, amount to call, and the complete action history. The model outputs an action in a structured format (e.g., check, bet with a specified size, call, or fold), which is sanitized to ensure compliance with game rules before execution.

This design isolates the effect of persona prompts on decision-making, allowing us to evaluate whether high-level behavioral descriptions alone can consistently steer an LLM toward coherent, strategically distinct, and human-interpretable poker behaviors over extended play.

4.4. Behavioral Feature Extraction

To quantify persona-induced behaviors, we simulate a fixed number of hands for each persona and compute standard PokerTracker-style statistics from the logged action histories. These statistics capture key strategic tendencies at different stages of play.

Specifically, we extract the following features: VPIP (voluntarily put money in pot), PFR (preflop raise frequency), aggression frequency on the flop, turn, and river ($AF_{q_{flop}}$, $AF_{q_{turn}}$, $AF_{q_{river}}$), WTSD (went to showdown), and W\$SD (won money at showdown). Together, these metrics summarize preflop looseness, postflop aggression, and showdown behavior.

Each opponent persona is represented as a feature vector,

$$\mathbf{x}_{\text{persona}} = [\text{VPIP}, \text{PFR}, \text{AF}_{q_{flop}}, \text{AF}_{q_{turn}}, \text{AF}_{q_{river}}, \text{WTSD}, \text{W\$SD}, \dots],$$

which serves as a compact empirical embedding of its strategic profile. These feature vectors form the basis for subsequent clustering and comparative analysis across personas.

4.5. Clustering Analysis

We examine whether LLM personas exhibit statistically separable behavioral patterns using both radar visualizations and low-dimensional projections of the extracted feature vectors.

Behavioral Radar Profiles. Radar charts constructed from PokerTracker-style statistics reveal clear and interpretable behavioral signatures across personas. Solver and TAG opponents exhibit balanced and stable aggression profiles, while the Nit persona shows extremely low VPIP and PFR, reflecting tight and risk-averse play. In contrast, the LAG persona displays consistently high aggression across all streets, and the Calling Station exhibits high showdown frequency combined with low aggression. These profiles closely align with widely accepted human poker archetypes.

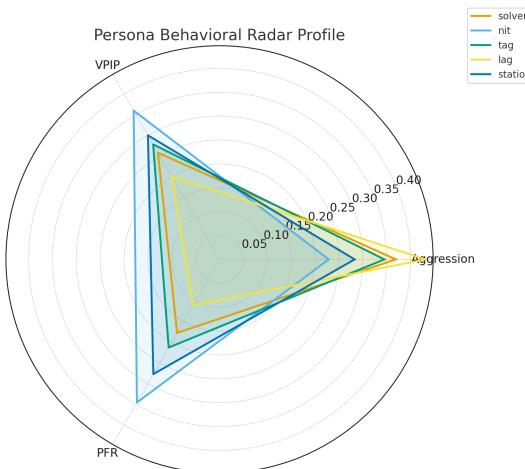


Figure 1: Behavioral radar profiles for different LLM personas based on aggregated poker statistics. Distinct aggression and passivity patterns emerge across personas.

PCA Projection. To further assess separability, we project the multi-dimensional behavioral feature vectors into a two-dimensional space using principal component analysis (PCA). The resulting projection shows clean clustering among personas: Nit and Calling Station occupy passive extremes, LAG is isolated along the high-aggression axis, and Solver and TAG lie near more balanced regions. This indicates that persona-conditioned LLM behaviors form coherent and distinct clusters in feature space.

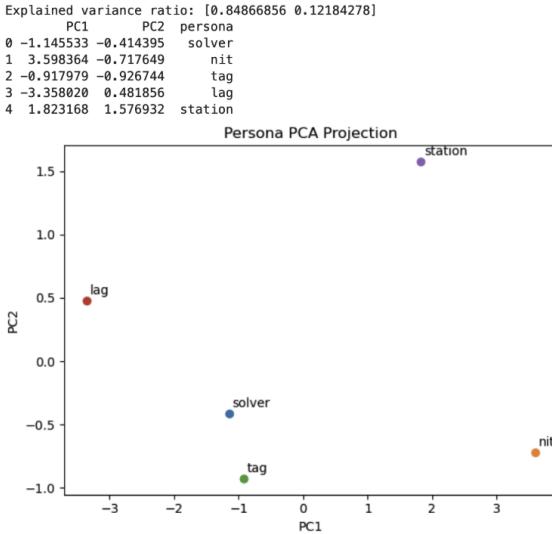


Figure 2: PCA projection of persona behavioral feature vectors. Personas form clearly separable clusters corresponding to different strategic styles.

Overall, these results suggest that LLM personas are not arbitrary labels but induce statistically and behaviorally meaningful differences consistent with real-world poker playing styles.

4.6. Performance Evaluation vs. Fixed Hero

We evaluate the impact of persona-conditioned opponent behavior by measuring game outcomes against a fixed, rule-based hero strategy. For each persona, the hero plays 1,000 independent hands under identical conditions, allowing differences in cumulative profit to be attributed primarily to opponent behavior.

The results reveal clear and interpretable performance differences across personas. The Nit persona collapses quickly, folding excessively and becoming highly exploitable. The Calling Station loses steadily over time, as frequent calling without sufficient aggression allows the hero to extract consistent value. The LAG persona exhibits high variance, achieving periods of strong performance but also experiencing large drawdowns due to over-bluffing. In contrast, the TAG and Solver personas perform competitively, yielding relatively stable outcomes and proving more difficult for the hero to exploit.

Overall, these outcome patterns align closely with theoretical expectations for real-world poker archetypes, providing further evidence that persona-conditioned LLM opponents exhibit meaningful and strategically coherent behavior.

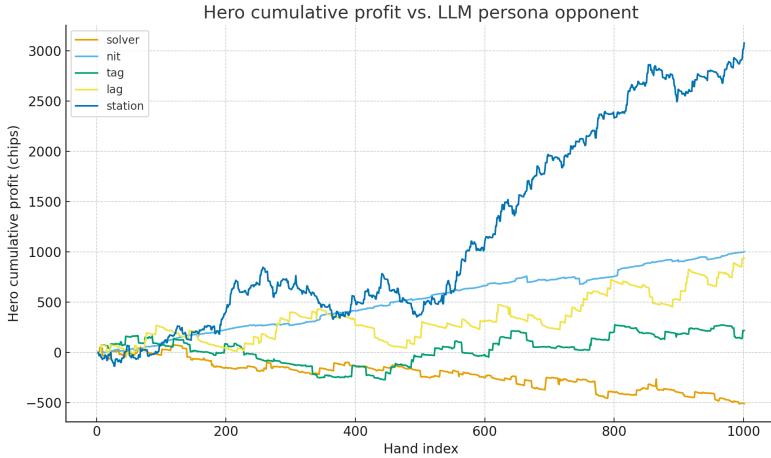


Figure 3: Hero cumulative profit over 1,000 hands against different LLM personas. Distinct performance trajectories reflect varying exploitability and strategic coherence across personas.

4.7. Key Findings

Across all experiments, persona-conditioned LLM opponents exhibit stable and reproducible strategic behaviors. Distinct personas produce clearly differentiated action-frequency profiles, leading to separable clusters in behavioral feature space and consistent differences in game outcomes against a fixed hero strategy. These patterns align with established human poker archetypes, particularly along aggression–passivity dimensions.

Overall, the results demonstrate that natural-language persona prompts provide a controllable and interpretable mechanism for inducing meaningful opponent styles in LLM-based agents, representing a practical step toward adaptive and behavior-aware opponent modeling.

5. Generative Opponent Model Embedding in RL Framework

5.1. From Behavior to Improvement

While the LLM-based study in Part 1 demonstrates that generative models can produce diverse and interpretable opponent behaviors, it does not directly address how such models can be used to improve learning and decision-making in reinforcement learning agents. In this section, we study a complementary setting in which a learned generative opponent model is embedded directly into an RL framework and trained jointly with the agent.

The key question motivating this part is: *Can we use a generative neural network as an opponent predictor to improve an RL poker agent’s performance?*

5.2. Environment and Framework Setup

We study whether explicit opponent modeling can improve reinforcement-learning-based poker agents in a simplified two-player Leduc Hold’em environment. Leduc Hold’em retains the essential imperfect-information and strategic structure of poker while dramatically reducing the state and action space. This simplification avoids the computational burden of full-scale games, enabling faster training and clearer analysis of learned behaviors.

The core idea is to embed opponent behavior prediction directly into the RL agent’s decision-making process. We train a generative neural model that predicts the opponent’s next action based on the current observable state. These predicted opponent actions are incorporated into the agent’s input features, allowing the policy to adapt dynamically based on anticipated opponent behavior.

We implement the environment using the PettingZoo interface, which supports turn-based multi-agent simulation. To isolate the effect of opponent modeling, we adopt a two-stage training framework:

- **Stage 1:** Train a baseline agent using Proximal Policy Optimization (PPO) against a random opponent.
- **Stage 2:** Freeze the Stage 1 policy as a fixed opponent. Train new PPO agents either with or without embedded opponent modeling to compare performance in controlled self-play.

This setup allows us to evaluate the impact of generative opponent modeling on policy learning in a controlled and reproducible environment, highlighting its benefits in terms of adaptability and win rate improvement.

5.3. Reinforcement Learning Agent via PPO

We train the poker agent using Proximal Policy Optimization (PPO), a policy-gradient method designed for stable on-policy reinforcement learning. PPO is well suited to our on-line poker setting, where trajectories are generated sequentially through interaction with the environment and the opponent.

State and Action Representation. The agent operates in a two-player Leduc Hold’em environment with a compact but expressive observation space. Each observation is a 36-dimensional vector encoding private cards, public cards, and chip counts for both players. Actions are discrete and correspond to the standard Leduc betting options: `call`, `raise`, `fold`, and `check`.

Table 1: Leduc Hold’em State Representation and Environment Specifications

| Index Range | Description |
|-------------|---|
| 0–2 | Current player’s private card (J, Q, K) |
| 3–5 | Public (community) card (J, Q, K) |
| 6–20 | Current player’s chip count (one-hot encoded) |
| 21–35 | Opponent’s chip count (one-hot encoded) |

| Property | Value |
|--------------------|--------------|
| Number of agents | 2 |
| Action space | Discrete (4) |
| Observation shape | (36,) |
| Observation values | [0, 1] |

Policy and Value Networks. The PPO agent consists of two neural networks:

- A policy network (actor) that maps the current state to action logits over the four legal actions.
- A value network (critic) that estimates the state value $V(s)$, which is used for advantage estimation.

Both networks are implemented as multilayer perceptrons (MLPs) and are trained jointly during policy optimization.

Advantage Estimation. To reduce variance while maintaining low bias, we use Generalized Advantage Estimation (GAE) with $\text{TD}(\lambda)$. Given rewards r_t and value estimates $V(s_t)$, the advantage is computed as

$$\hat{A}_t = \sum_{l=0}^{\infty} (\gamma\lambda)^l \delta_{t+l}, \quad \delta_t = r_t + \gamma V(s_{t+1}) - V(s_t),$$

where γ is the discount factor and λ controls the bias–variance tradeoff.

PPO Objective. PPO updates the policy by maximizing a clipped surrogate objective:

$$\mathcal{L}^{\text{PPO}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right],$$

where

$$r_t(\theta) = \frac{\pi_\theta(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)},$$

and ϵ is the clipping parameter that constrains policy updates.

Key Properties. This PPO-based design provides several advantages for our poker setting:

- Stable policy updates through clipped optimization.
- On-policy learning, which aligns naturally with online self-play and opponent interaction.
- Compatibility with opponent modeling, as state representations can be augmented with predicted opponent behavior without changing the optimization framework.

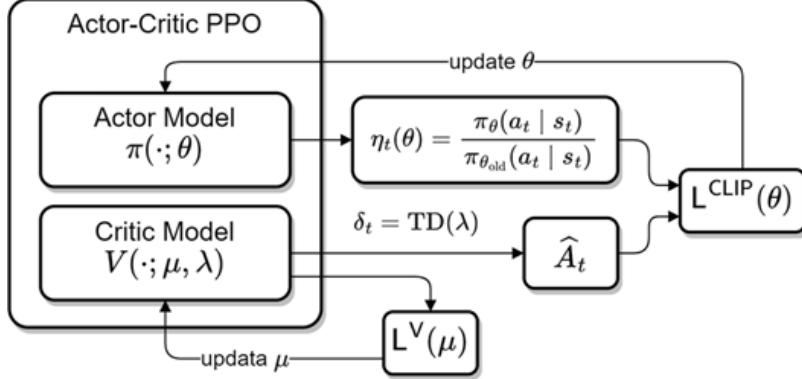


Figure 4: PPO architecture

5.4. Opponent Modeling via MLP

To enhance policy adaptation, we incorporate an explicit opponent prediction model that estimates the opponent’s next action from the current state. Unlike persona-conditioned LLMs used in Part 1, this model is fully data-driven and integrated into the reinforcement learning framework.

Opponent Predictor Network. The opponent model is implemented as a lightweight multilayer perceptron (MLP) classifier. It receives the same 36-dimensional observation vector as the agent, which encodes the current hand, public card, and chip counts. The model outputs a probability distribution over the opponent’s four possible actions: `call`, `raise`, `fold`, and `check`.

Architecture Details. The network has two hidden layers with ReLU activations, each with 32 neurons, followed by a softmax output layer. The architecture is simple but effective enough in capturing short-term behavioral signals.

- **Input:** Current 36-dimensional observation
- **Hidden Layers:** Two fully connected layers with ReLU (each 32 units)
- **Output:** Softmax over 4 opponent action classes

Training is performed online during PPO agent learning using cross-entropy loss and the Adam optimizer. Each time the opponent acts, the corresponding state-action pair is added as a supervised learning example for the model.

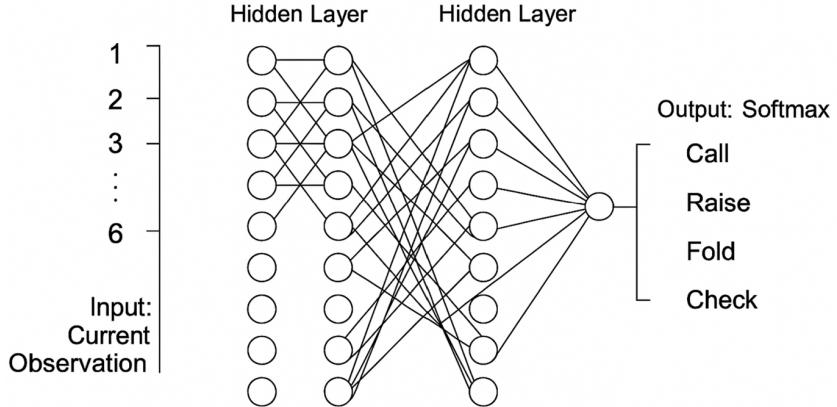


Figure 5: Opponent predictor neural network architecture.

Integration into RL Agent. To incorporate predicted opponent behavior into decision-making, the PPO agent augments its input state with the predicted opponent action probabilities. Specifically, the 4-dimensional output of the opponent model is concatenated with the raw observation vector, yielding a 40-dimensional input to the policy and value networks.

This setup enables the PPO policy to condition its actions not just on the observed state, but also on a learned model of the opponent’s likely behavior. The joint training setup reflects a live gameplay environment, where both the policy and its internal belief about the opponent evolve together.

This integrated architecture allows the agent to dynamically adapt to different opponent tendencies and offers a foundation for more behavior-aware reinforcement learning in imperfect-information games like poker.

5.5. Training Settings

We design a two-stage training framework to isolate the effect of explicit opponent modeling within a reinforcement learning pipeline and enable controlled comparisons.

Stage 1: Baseline Training. In the first stage, a PPO agent is trained against an opponent that follows a uniformly random strategy. This setting exposes the agent to a wide variety of game situations while avoiding strong strategic bias from the opponent. The objective of this stage is to learn a reasonable but naive baseline poker policy that captures fundamental betting and folding behavior. After convergence, the trained PPO policy is frozen and used as a fixed opponent in subsequent experiments.

Stage 2: Evaluation and Comparison. In the second stage, we evaluate whether opponent modeling improves learning and performance. The frozen PPO agent obtained from Stage 1 serves as a fixed opponent, ensuring a stationary and controlled evaluation environment. Two agents are trained and compared under identical conditions:

- A PPO agent without opponent modeling, which observes only the raw environment state.

- A PPO agent with opponent modeling, which receives an augmented state that includes predicted opponent action probabilities.

This setup allows performance differences to be attributed directly to the inclusion of opponent modeling.

Opponent Model Training. The opponent predictor is trained online and synchronously during gameplay. At each opponent decision step, the observed state and the opponent’s executed action are treated as a supervised learning example. One gradient update is performed per opponent action, mimicking a live, non-stationary gameplay scenario in which the agent continuously refines its beliefs about opponent behavior as new evidence becomes available.

Key Hyperparameters. Across all experiments, we use a consistent set of hyperparameters to ensure comparability:

- PPO training iterations: 30
- Episodes per iteration: 1000
- Advantage estimation: $\text{TD}(\lambda)$ with $\lambda = 0.9$
- PPO clipping parameter: $\epsilon = 0.2$
- Optimizer: Adam for both policy and value networks

These settings balance training stability and computational efficiency while remaining suitable for the simplified Leduc Hold’em environment.

5.6. Performance Evaluation

We evaluate the performance of PPO agents with and without opponent modeling across two training stages, and analyze the quality of the opponent model itself.

Stage 1: Baseline PPO Training

In the first stage, the PPO agent is trained against a random opponent. As shown in Figure 6, the agent’s average reward increases steadily and the win rate quickly rises to around 90%. This indicates successful convergence to a competent baseline policy.

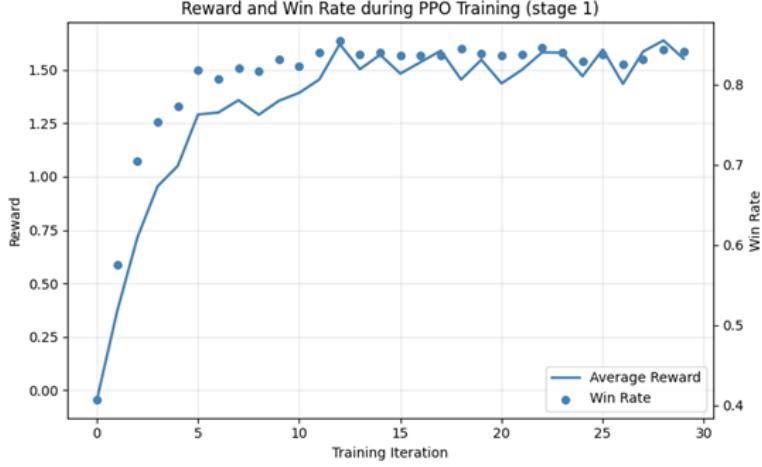


Figure 6: Reward and win rate during PPO training against random opponent (Stage 1).

Opponent Model Learning Quality

The opponent predictor network learns rapidly during training. As shown in Figure 7, the training loss decreases sharply, while prediction accuracy reaches and maintains around 90% after approximate 2,000 steps. This confirms that opponent behavior is learnable from raw observations, even under online, non-stationary conditions.

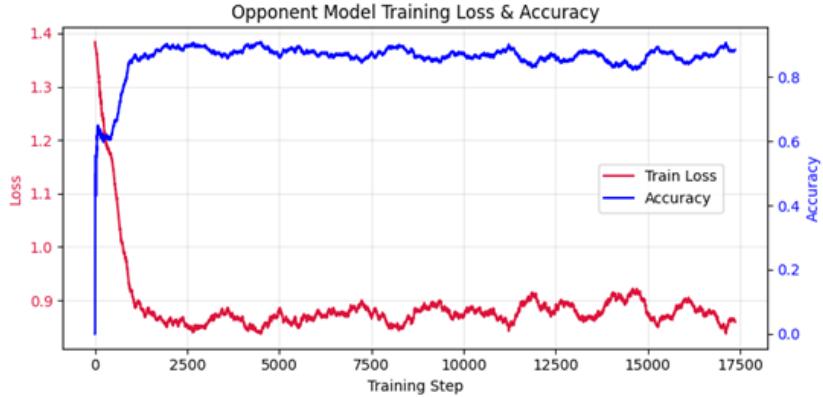


Figure 7: Opponent model training loss and accuracy over time.

Stage 2: Agent Comparison with vs. without Opponent Modeling

In the second stage, two PPO agents are trained against the fixed policy from Stage 1. As shown in Figure 8, both agents show improved reward and win rate over time.

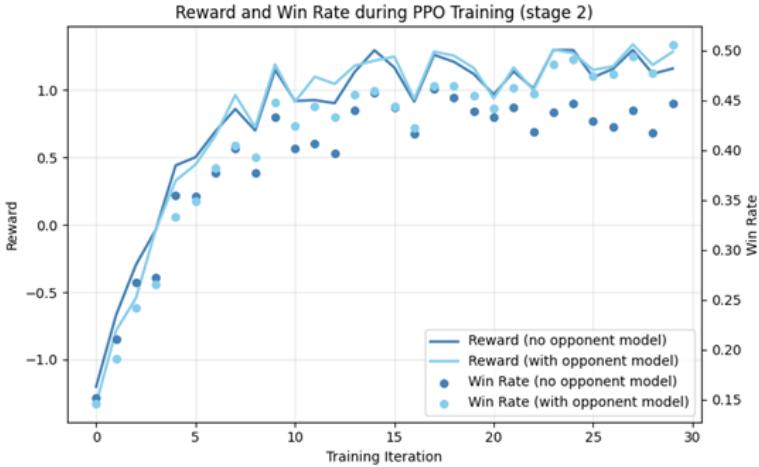


Figure 8: Comparison of reward and win rate between agents with and without opponent model (Stage 2).

Final Metrics

At the end of training:

- Final win rate (no opponent model): 44.7%
- Final win rate (with opponent model): 50.6%
- Improvement over baseline: +5.9 percentage points ($\approx 13.2\%$ relative gain)

These results validate the effectiveness of embedding generative opponent models in reinforcement learning agents. Even a simple predictive model provides actionable behavior signals that improve policy adaptation and overall performance in imperfect-information games.

6. Future Works and Potential Improvements

Several extensions can further enhance the proposed LLM-based opponent modeling framework. First, opponent conditioning can be made richer and more flexible by learning persona representations directly from gameplay data. Instead of relying on fixed natural-language prompts, future work could learn continuous persona embeddings that capture a broader spectrum of strategic behaviors and allow smoother interpolation between opponent styles.

Second, the current approach assumes that opponent behavior remains stationary throughout play. An important direction is online adaptation, where the opponent model updates its representation dynamically as new actions are observed. This would enable agents to track evolving strategies and better reflect realistic human opponents.

Finally, the framework can be scaled to more complex environments. Extending the evaluation to multi-player poker settings and longer decision horizons would provide a more comprehensive assessment of the effectiveness and robustness of LLM-based opponent modeling in strategic games.

In parallel, several enhancements can be made to the reinforcement learning framework with embedded opponent prediction. More nuanced training settings—such as experience replay, curriculum learning over opponent skill levels, and asynchronous or delayed updates—can improve training stability and mitigate feedback loops between the policy and the opponent model.

Moreover, richer representations of opponent behavior could be achieved by modeling long-term patterns such as bluff frequency, aggression sequences, or street-dependent tendencies. Recurrent architectures or trajectory-level embeddings may provide a more expressive and temporally coherent understanding of opponent styles.

Finally, adapting this framework to more complex poker variants—including multi-player Leduc or full-scale Texas Hold’em—requires additional components such as public belief modeling and deeper action abstraction layers. These improvements would help handle the increased uncertainty and strategic diversity in larger, more realistic game settings.

7. Conclusion

This work presents a unified study of generative opponent modeling in poker. In the first part, we show that persona-conditioned large language models can reliably generate distinct, human-interpretable opponent behaviors that align with established poker archetypes. These behaviors are statistically separable and produce meaningful differences in game outcomes.

In the second part, we demonstrate that generative opponent modeling can be embedded directly into a reinforcement learning framework. By conditioning a PPO agent on a learned opponent action predictor, the agent achieves higher rewards and win rates, illustrating how opponent representations can actively inform decision-making.

Together, these results suggest a coherent framework in which generative models serve dual roles: as simulators of strategic diversity and as structured embeddings that guide adaptive policies. This connection highlights the potential of generative AI not only to model opponents, but to fundamentally reshape how agents learn and reason in imperfect-information games.

References

- Brown, N. and Sandholm, T. (2018). Superhuman ai for heads-up no-limit poker: Libratus beats top professionals. *Science*.
- Chen, L., Lu, K., Rajeswaran, A., and et al. (2021). Decision transformer: Reinforcement learning via sequence modeling. *NeurIPS*.
- Ekmekci, O. and Sirin, V. (2013). Learning strategies for opponent modeling in poker. In *Computer Poker and Imperfect Information: Papers from the AAAI 2013 Workshop*, pages 6–12.
- Goodfellow, I. et al. (2014). Generative adversarial nets. *NeurIPS*.
- Heinrich, J. and Silver, D. (2016). Deep reinforcement learning from self-play in imperfect-information games. *arXiv preprint arXiv:1603.01121*.
- Ho, J. and Ermon, S. (2016). Generative adversarial imitation learning. In *NeurIPS*.

- Ho, J. and Salimans, T. (2020). Denoising diffusion probabilistic models. *NeurIPS*.
- Project, S. U. C. (2024). Llm-guided strategy and opponent modeling in multi-agent poker.
- Vaswani, A. et al. (2017). Attention is all you need. In *NeurIPS*.